



Bank Loan Approval Prediction

Credit Risk Assessment, Performance Diagnostics & Production
Engineering Roadmap



Business Context & Strategic Objectives

Modernizing Credit Underwriting: Transition from manual assessments to automated, data-driven scoring.

Key Objectives

- Accelerate decision turnaround times to enhance operational throughput.
- Accurately screen out high-risk applicants to minimize default exposure.
- Build a transparent, regulatory-compliant machine learning pipeline for governance.

Core Business Challenge

Balancing **rapid loan approval throughput** against **default exposure** and credit losses.

Automated ML-driven scoring provides the required precision to manage this asymmetric risk profile.



Dataset Architecture & Data Health

TOTAL APPLICATIONS

614

CANDIDATE FEATURES

12

DATA HEALTH ALERT

- 149 missing values identified across financial/demographics.
- Skewed distribution: 68.7% Approved / 31.3% Rejected.

DOMAIN GROUP	FEATURE MAPPING & ATTRIBUTES
Financial	ApplicantIncome, CoapplicantIncome, LoanAmount, Loan_Amount_Term
Credit History	Credit_History (Guidelines Compliance vs. No History)
Demographics	Gender, Married, Dependents, Education, Self_Employed, Property_Area



Baseline Model Architecture

Total Applications (Test)

123

Model Type

Random Forest

Test Accuracy

78.05%

Pipeline Configuration

- Global LabelEncoder utilized for categorical feature preprocessing.
- Dataset split into **80% Train / 20% Test** (96/123 correct classifications).
- Random Forest Classifier ensemble constructed with **100 Trees**.



Performance Diagnostics & Risk Asymmetry

CLASS 0 RECALL (REJECTED)

42%

CLASS 1 RECALL (APPROVED)

97%

RISKY MISS RATE

58.1%

Confusion Matrix (N=123)

18

Correct Rejections

25

False Approvals

2

False Rejections

78

Correct Approvals

Diagnostic Findings

- **Class 0 (Rejected):** High Precision (90%) but low Recall (42%), indicating conservative risk detection.
- **Class 1 (Approved):** Robust performance with 97% Recall and 0.85 F1-score for standard applicants.
- **Critical Risk:** 58.1% miss rate on risky applicants leads to 25 false approvals, creating high loss potential.
- **Model Asymmetry:** Performance is heavily skewed toward approval volume over risk mitigation objectives.



Financial Risk Analysis: Cost of Errors

Asymmetric Risk in Banking

False Approvals (Critical Severity)

25 Cases

Leads to loan default, capital loss, recovery and legal costs.

False Rejections (Low Severity)

2 Cases

Leads to minor friction and lost interest income.

Underwriting Imperative: Optimize models to penalize False Approvals.



Technical Limitations of Baseline Pipeline



Critical Pipeline Flaws

Identifier Leakage: `Loan_ID` included in training, causing spurious correlation.


Preprocessing Leakage: `LabelEncoder` fit globally before split, introducing future data knowledge.

Artificial Ordinality: Nominal features (`Property_Area`, `Gender`) improperly encoded as integers.

Unhandled Missing Values: 149 data points not systematically imputed, leading to record exclusion or bias.



Proposed Production ML Pipeline



Stage 1: Pre-Modeling Pipeline

- **Preprocessing**: Automated `ColumnTransformer` with Median Imputation for missing values, Robust Scaling for outliers, and One-Hot Encoding.
- **Feature Engineering**: High-impact domain features including Total Income, DTI Ratio, and EMI-to-Income indicators.

Stage 2: Ensemble & Optimization

- **Optimization**: Class imbalance mitigation via SMOTE and Balanced Weights; Threshold Tuning calibrated for a **80% recall target**.
- **Models**: Multi-model ensemble consisting of XGBoost, LightGBM, and Cost-Weighted Logistic Regression for robust inference.



Model Governance & Fair Lending Compliance

Regulatory Framework

- Regulatory Compliance: Strict adherence to [ECOA](#) and [FCRA](#) requirements.
- Audit demographic variables to eliminate disparate impact across protected classes.
- Ensuring rigorous data lineage and feature selection transparency.

Explainability Engine

- Deployment of [SHAP](#) and [LIME](#) for model interpretability.
- **Global:** Ensure aggregate behavior aligns with established credit theory.
- **Local:** Automated generation of Adverse Action Reason Codes for every rejection.

Transparent AI: Driving organizational trust by balancing algorithmic complexity with absolute regulatory accountability and fair lending standards.



Implementation Roadmap & Milestones

PHASE 1: CODE REMEDIATION

Timeline: Immediate

- Clean pipeline architecture
- Drop Loan_ID & One-Hot Encode
- Robust Median Imputation

PHASE 2: BENCHMARKING

Timeline: Weeks 1-2

- Train XGBoost & LightGBM
- Threshold tuning for risk
- Target $\geq 80\%$ detection rate

PHASE 3: DEPLOYMENT

Timeline: Weeks 3-4

- Integrate SHAP microservice
- CI/CD API endpoint production
- Real-time monitoring drift